

# Sort'em

*Analyzing various sorting algorithms in three different programming languages*

Adam Bowker  
Sai Macha  
ST 370-601

March 28, 2018

# Executive Summary

Our goal with this project was to find the fastest and most efficient method of sorting a list of objects, based on sorting algorithm and programming language. Our response variable is the time taken to sort the list, and the factors are **input size** (100, 1000, and 10000), **sorting algorithm** (MergeSort, BubbleSort, SelectionSort, and QuickSort), and **programming language** (Java, C, and Python). We used the *completely randomized design* method, so each trial was run based on mapping our design matrix to a random array of integers from 0 to 36. We also randomized in the creation of our lists to sort. First, given the input size, we generated a list of integers 0 up to (size - 1). We then used random number generators from each of the languages' default packages, seeded with the system's current time, and shuffled the list. This list was then copied 4 times, and sorted/timed on its own by each of the algorithms.

In the end, we discovered QuickSort to be the fastest sorting algorithm, regardless of input size or chosen language. As far as language efficiency goes, Java preformed the best out all three languages, given an algorithm and an input size. Overall, QuickSort in Java seems to be the most efficient way to sort a list of objects, regardless of input size.

Each of these algorithms is very popular and widely used for a variety of purposes. QuickSort is recognized as the lowest-cost (fastest) sorting algorithm, but sometimes a slower one (i.e. BubbleSort, known for being incredibly costly) is used because efficiency is not an issue and implementation is easier. For example, QuickSort may be implemented by a team of developers working on a large-scale project where efficiency is crucial, but a student working on a small project for sorting relatively small lists may choose to implement BubbleSort because it's only a few lines of easy-to-understand code. For the purposes of this study, each algorithm was implemented using the same data structures (standard, statically-allocated arrays) and object types (Integers).

## Table of Contents

Executive Summary	2
Reason for the Study	4
Prior Expectations	5
What We Did	6
Raw Data	9
Statistical Analysis	10
Practical Implications	22
Further Questions	23
Original Project Proposal	(separate file)

## Reason for the Study

We chose this study because of the wide implications of the results. Understanding what goes into sorting lists and how efficient the various sorting algorithms are is important for anyone who is developing anything, from video games to large-scale “backend” software to even a basic “Hello World”-type script. For example, for my CSC 316 class, I was working on a project where efficiency was critical, and I needed to have this kind of information so that I could choose the best sorting algorithm. Another example that can be applied to statistics: a certain set of data needs to be sorted before it can be analyzed, what’s the best way to do it? Of course, this sorting will likely only happen once, so efficiency may not be crucial. However, if you take a look at our raw data on GitHub ([data.txt](#)), the same set of data (100,000 integers) took .015 seconds with QuickSort in C, but over an hour using BubbleSort in Python - sometimes, an hour can be pretty unreasonable when you’ve just had a statistical breakthrough and need to get the data sorted fast!

## Prior Expectations

Before we started data collection, we weren't entirely sure what to expect. We had used sorting algorithms in class before, but never extensively enough to know exactly how MergeSort in Java and QuickSort in Python (for example) would compare. However, we did have enough information about the 3 languages and 4 algorithms to make a fairly educated guess. Java and C are both general purpose programming languages (and need to be compiled), whereas Python is typically recognized as a scripting language, and does not require compilation beforehand.

In addition to this, each of the 4 algorithms has previously been analyzed at a high level, and their best-case, worst-case, and average-case runtimes are proven. They are as follows:

Algorithm	Best-case	Worst-case	Average-case
BubbleSort	$O(n)$	$O(n^2)$	$O(n^2)$
MergeSort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$
QuickSort	$O(n \log n)$	$O(n^2)$	$O(n \log n)$
SelectionSort	$O(n^2)$	$O(n^2)$	$O(n^2)$

Read more about Big-O notation here: <https://rob-bell.net/2009/06/a-beginners-guide-to-big-o-notation/>

This information allowed to make a prediction that QuickSort would be the fastest, and we predicted that it would be fastest in C as it is not object-oriented like Python and Java. Another prediction could have been that QS and/or MS would be fastest in Python, because it is interpreted rather than compiled.

Our predictions were somewhat correct; the fastest algorithm was definitely QuickSort, regardless of language or input size. The fastest language was surprising though - it was Java, which was unexpected as it needs to be compiled and needs to run on a virtual machine.

# What We Did

In order to test these algorithms in the different languages, first we had to write the code. We wrote three programs, one in Java, one in C, and one in Python, each of which does the following:

- Read an input size from the command line
- Generate a list of random, unique integers from 0 to the input size - 1
- Copy the list 4 times, once for each algorithm
- Sort each copy of the list using one of the four algorithms, and time how long it took
- Output the data in an easy-to-consume format

One concern going into the project was that our code would not necessarily be the most efficient, or that the efficiency would not be consistent across languages. To combat this, we used the exact same data structure in each algorithm (standard statically-allocated array), and only began the clock for timing the algorithm after the random lists had already been generated (in some cases, such as QuickSort in C, the time taken to generate the random list was actually longer than the time it took to sort it!) and before the results were printed, so that the recorded time *only* took into account the exact time taken to sort the list. These algorithms are also very well-known and well-established, so we implemented them exactly as they came, with no modifications that could alter efficiency.

Pseudo-code for each algorithm can be found below, or our full code can be found on GitHub: <https://github.com/adboio/st370>. Instructions for compiling and running the code can be found within the README.md file in the linked repository, but for the ease of reading, they are copied below, along with sample output:

*\*The commands are to be used in Terminal on Mac, PuTTY on PC, or Command Line on Linux. These all assume you have previously navigated to the location of the source file and have the appropriate programming language installed. No special packages are required.*

## Algorithms' Pseudo-code

### BubbleSort - $O(n^2)$

```
bubbleSort(list):  
    for all elements in list  
        if list[i] > list[i+1] then  
            swap list[i] and list[i+1]
```

### SelectionSort - $O(n^2)$

```
selectionSort(list):  
    for i = 0 up to size of list  
        min = i  
        for j = i+1 up to size of list  
            if list[j] < list[min]  
                Min = j  
        if min does NOT equal i then  
            swap list[min] and list[i]
```

### MergeSort - $O(n \log n)$

```
mergeSort(list list):  
    if length of list <= 1 then  
        return list  
    right = empty list  
    left = empty list  
    for each element with index i in list do  
        if i < (length of list)/2 then  
            add element to left  
        else  
            add element to right  
    right = mergeSort(right)  
    left = mergeSort(left)  
    return merge(left, right)  
  
merge(list right, list left):  
    result = empty list  
  
    while left NOT empty and right NOT empty do  
        if first element in left <= first element in right then  
            add first element in left to result  
            remove first element from left  
        else  
            add first element in right to result  
            remove first element from right  
  
    while left is NOT empty do  
        add first element in left to result  
        remove first element from left  
    while right is NOT empty do  
        add first element in right to result  
        remove first element from right  
  
    return result
```

## QuickSort - $O(n \log n)$ :

```
quickSort(list, low, high):
    If low < high then
        Pi = partition(list, low, high)

        quickSort(list, low, pi-1)
        quickSort(list, pi+1, high)

partition(list, low, high):
    pivot = list[high]
    i = low - 1
    for each element j in list up to index (high - 1) do
        if list[j] <= pivot then
            i = i+1
            swap arr[i] and arr[j]
    swap list[i + 1] and list[high]
    return i + 1
```

## How to Run (+ sample output)

### sorting.c

compile: `gcc -Wall -std=c99 sorting.c -o sorting`

run: `./sorting <size>`

example (with output):

`./sorting 100`

~~~~~RESULTS~~~~~

| <Algorithm>   | <Runtime> |
|---------------|-----------|
| MergeSort     | 0.000015  |
| QuickSort     | 0.000011  |
| SelectionSort | 0.000020  |
| BubbleSort    | 0.000037  |

### ST370.java

compile: `javac ST370.java`

run: `java ST370 <size>`

example (with output):

`java ST370 100`

~~~~~RESULTS~~~~~

| <Algorithm>   | <Runtime> |
|---------------|-----------|
| MergeSort     | 729241    |
| QuickSort     | 283261    |
| SelectionSort | 74466     |
| BubbleSort    | 129271    |

### sorting.py

(no compilation needed)

run: `python sorting.py <size>`

example (with output):

`python sorting.py 100`

~~~~~RESULTS~~~~~

| <Algorithm>    | <Runtime>         |
|----------------|-------------------|
| MergeSort:     | 0.000759840011597 |
| QuickSort:     | 0.000298023223877 |
| SelectionSort: | 0.000952959060669 |
| BubbleSort:    | 0.0018670558929   |



## Raw Data

Our raw data can be found in our repository ([github.com/adboio/st370](https://github.com/adboio/st370)), or pasted below:

| Language | Input Size | Algorithm     | Time        | Trial |
|----------|------------|---------------|-------------|-------|
| Java     | 100        | MergeSort     | 0.001836882 | 31    |
| Java     | 100        | QuickSort     | 0.000863929 | 12    |
| Java     | 100        | SelectionSort | 0.000087507 | 24    |
| Java     | 100        | BubbleSort    | 0.000148425 | 5     |
| Java     | 1000       | MergeSort     | 0.00240215  | 22    |
| Java     | 1000       | QuickSort     | 0.001140689 | 27    |
| Java     | 1000       | SelectionSort | 0.002719119 | 8     |
| Java     | 1000       | BubbleSort    | 0.003994092 | 26    |
| Java     | 10000      | MergeSort     | 0.003588684 | 19    |
| Java     | 10000      | QuickSort     | 0.002564156 | 15    |
| Java     | 10000      | SelectionSort | 0.055551656 | 13    |
| Java     | 10000      | BubbleSort    | 0.12884229  | 21    |
| C        | 100        | MergeSort     | 0.000016    | 18    |
| C        | 100        | QuickSort     | 0.000011    | 32    |
| C        | 100        | SelectionSort | 0.000021    | 30    |
| C        | 100        | BubbleSort    | 0.000038    | 9     |
| C        | 1000       | MergeSort     | 0.000166    | 28    |
| C        | 1000       | QuickSort     | 0.000107    | 33    |
| C        | 1000       | SelectionSort | 0.001299    | 16    |
| C        | 1000       | BubbleSort    | 0.002906    | 14    |
| C        | 10000      | MergeSort     | 0.001467    | 7     |
| C        | 10000      | QuickSort     | 0.001156    | 36    |
| C        | 10000      | SelectionSort | 0.090376    | 4     |
| C        | 10000      | BubbleSort    | 0.283806    | 20    |
| Python   | 100        | MergeSort     | 0.001095057 | 1     |
| Python   | 100        | QuickSort     | 0.000466108 | 34    |
| Python   | 100        | SelectionSort | 0.001313925 | 6     |
| Python   | 100        | BubbleSort    | 0.002349138 | 2     |
| Python   | 1000       | MergeSort     | 0.010231018 | 35    |
| Python   | 1000       | QuickSort     | 0.004241943 | 10    |
| Python   | 1000       | SelectionSort | 0.088142872 | 23    |
| Python   | 1000       | BubbleSort    | 0.190106154 | 25    |
| Python   | 10000      | MergeSort     | 0.128978014 | 3     |
| Python   | 10000      | QuickSort     | 0.054528952 | 29    |
| Python   | 10000      | SelectionSort | 9.394102097 | 17    |
| Python   | 10000      | BubbleSort    | 19.8599689  | 11    |

# Statistical Analysis

We used the stats function to analyze the data split by programming language, size of calculation, and algorithm used to sort.

```
>> stats(sort.time, sort.language)
```

|           | Java        | C           | Python      |
|-----------|-------------|-------------|-------------|
| N         | 12.00000000 | 12.00000000 | 12.00000000 |
| Mean      | 0.01697800  | 0.0317810   | 2.47800000  |
| Std. Dev. | 0.03847200  | 0.0834520   | 6.09800000  |
| Q1        | 0.00093312  | 2.525e-05   | 0.00157270  |
| Median    | 0.00248320  | 0.0006610   | 0.03238000  |
| Q3        | 0.00389270  | 0.0025462   | 0.17482000  |
| Min       | 8.7507e-05  | 1.1e-05     | 0.00046611  |
| Max       | 0.12884000  | 0.2838100   | 19.86000000 |
| Range     | 0.12875000  | 0.2838000   | 19.85950000 |

From the stats output for programming language it can be deduced that python takes the longest on average with a mean of 2.47s and the largest standard deviation.

Python had the largest maximum time to compute, much higher than the mean, which attributes to the large standard deviation. Java is the quickest on average at 0.0169s; however, the language C had the lowest minimum at 1.1e-5s.

```
>> stats(sort.time, log10(sort.size))
```

|           | 2           | 3           | 4           |
|-----------|-------------|-------------|-------------|
| N         | 12.00000000 | 12.00000000 | 12.00000000 |
| Mean      | 0.00068725  | 0.0256210   | 2.5004000   |
| Std. Dev. | 0.00080600  | 0.0573630   | 6.0885000   |
| Q1        | 2.525e-05   | 0.0011803   | 0.0028203   |
| Median    | 0.00030727  | 0.0028126   | 0.0729640   |
| Q3        | 0.00125920  | 0.0087337   | 0.2451000   |
| Min       | 1.1e-05     | 0.0001070   | 0.0011560   |
| Max       | 0.00234910  | 0.1901100   | 19.8600000  |
| Range     | 0.00233810  | 0.1900000   | 19.8588000  |

For the stats output for size, we had to analyze the data with a logarithmic function.

We did this throughout the statistical analysis because the variations were too high. As

shown in the output, 2 (  $\log_{10}(100)=2$  ) has the smallest standard deviation and 4 has the highest; 4 also has the highest average with 2 having the lowest. This indicates that as the size of the sorting gets larger, the longer it takes.

```
>> stats(sort.time, sort.algorithm)
```

|                  | MergeSort  | QuickSort  | SelectionSort | BubbleSort |
|------------------|------------|------------|---------------|------------|
| <b>N</b>         | 9.00000000 | 9.00000000 | 9.00000000    | 9.00000000 |
| <b>Mean</b>      | 0.01664200 | 0.00723110 | 1.07040000    | 2.2747000  |
| <b>Std. Dev.</b> | 0.04223800 | 0.01778700 | 3.12160000    | 6.5953000  |
| <b>Q1</b>        | 0.00063053 | 0.00028655 | 0.00069325    | 0.0012488  |
| <b>Median</b>    | 0.00183690 | 0.00114070 | 0.00271910    | 0.0039941  |
| <b>Q3</b>        | 0.00690990 | 0.00340300 | 0.08925900    | 0.2369600  |
| <b>Min</b>       | 1.6e-05    | 1.1e-05    | 2.1e-05       | 3.8e-05    |
| <b>Max</b>       | 0.12898000 | 0.05452900 | 9.39410000    | 19.8600000 |
| <b>Range</b>     | 0.12896000 | 0.05451800 | 9.39410000    | 19.8599000 |

In this output we analyzed the sorting algorithm. QuickSort had the smallest mean and bubble sort had the largest. Quicksort had the smallest standard deviation and smallest range, indicating it was more efficient than the other sorting algorithms when the sizes got larger. Bubblesort was the slowest to sort the data out of all of the sorting algorithms with the largest maximum and minimum.

## Mfit

```
>> mfit(sort.time, sort.language, log10(sort.size), sort.algorithm)
```

Overall Mean  
0.84224

Fitted Main Effect of Y variable , y, by X variable, x1

| Source | N  | Main Effect |
|--------|----|-------------|
| Java   | 12 | -0.82526    |
| C      | 12 | -0.81046    |
| Python | 12 | 1.63570     |

Fitted Main Effect of Y variable , y, by X variable, x2

| Source | N  | Main Effect |
|--------|----|-------------|
| 2      | 12 | -0.84155    |
| 3      | 12 | -0.81662    |

4      12      1.65820

Fitted Main Effect of Y variable , y, by X variable, x3

| Source        | N | Main Effect |
|---------------|---|-------------|
| MergeSort     | 9 | -0.82560    |
| QuickSort     | 9 | -0.83501    |
| SelectionSort | 9 | 0.22816     |
| BubbleSort    | 9 | 1.43240     |

Table of 2-way x1 by x2 Interaction Effects

|    |   | x1       |          |         |
|----|---|----------|----------|---------|
|    |   | Java     | C        | Python  |
| x2 | 2 | 0.82531  | 0.80979  | -1.6351 |
|    | 3 | 0.80220  | 0.78596  | -1.5882 |
|    | 4 | -1.62750 | -1.59580 | 3.2233  |

Table of 2-way x1 by x3 Interaction Effects

|    |               | x1       |          |          |
|----|---------------|----------|----------|----------|
|    |               | Java     | C        | Python   |
| x3 | MergeSort     | 0.81123  | 0.79437  | -1.60560 |
|    | QuickSort     | 0.81955  | 0.80365  | -1.62320 |
|    | SelectionSort | -0.22569 | -0.22938 | 0.45506  |
|    | BubbleSort    | -1.40510 | -1.36860 | 2.77370  |

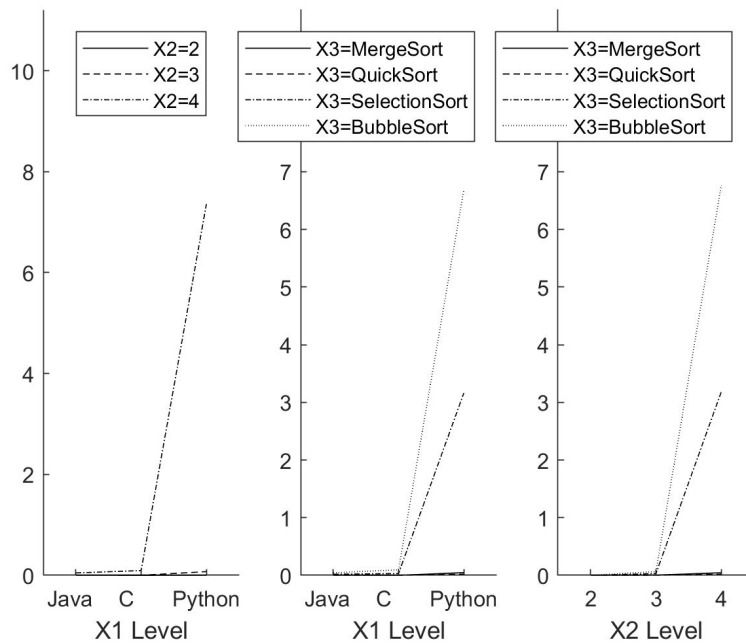
Table of 2-way x2 by x3 Interaction Effects

|    |               | x2       |          |          |
|----|---------------|----------|----------|----------|
|    |               | 2        | 3        | 4        |
| x3 | MergeSort     | 0.82589  | 0.80424  | -1.63010 |
|    | QuickSort     | 0.83477  | 0.81122  | -1.64600 |
|    | SelectionSort | -0.22837 | -0.22306 | 0.45144  |
|    | BubbleSort    | -1.43230 | -1.39240 | 2.82470  |

\*With absolute values of 1.6370, 1.65820, and 1.4320 for language(x1), size(x2), and algorithm(x3), respectively, we can assume that the size has the largest impact on the speed of sorting. The largest absolute value indicates which of the variables has the largest impact on the sort speed. The table for 2-way interactions also give us an idea of which factors have a larger role.

## Mplot

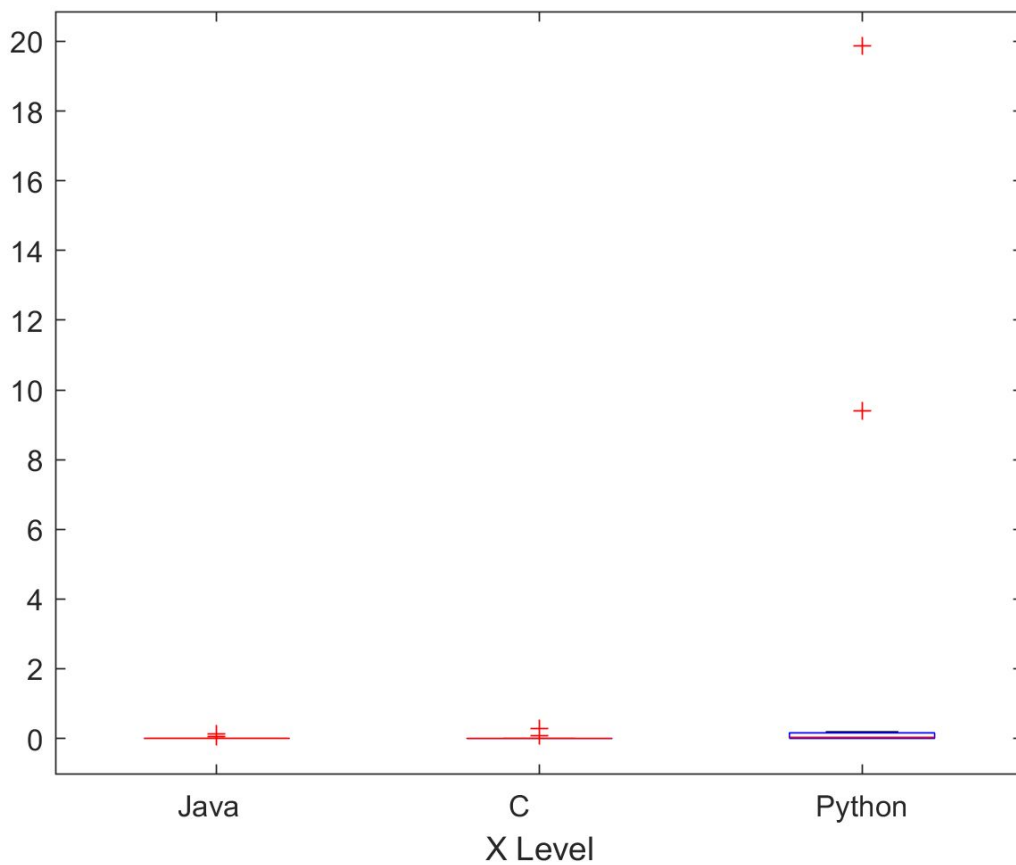
```
>> mplot2(sort.time, sort.language, log10(sort.size), sort.algorithm)
```



A mplot graph shows us the interaction, or lack thereof, between different levels. In this output of mplot2, There are indications of interaction between the different levels. Since the slopes are different, we know that there is interaction between the levels. The positive slopes in all the graphs indicate that the next level has a greater mean than the level prior; for example, C has a higher mean than Java with size 4. The lines start out close together and then spread out rapidly indicating that the difference is not as important in the beginning, but increases rapidly when moving to different levels.


# Bplot

```
>> bplot(sort.time, sort.language)
```



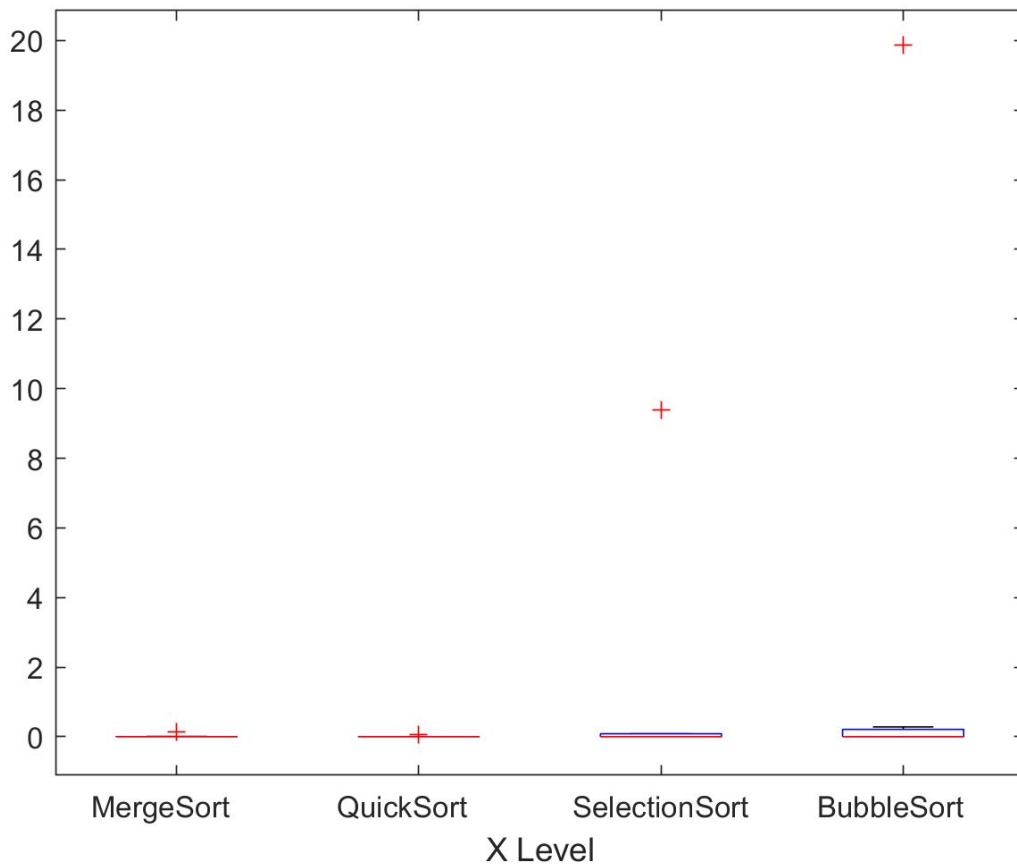
From this bplot we can see that the range is far greater in python than the other two computer languages. We can infer that the standard deviation is also much higher in python because there is a box visible while the other two languages are skewed into a line.

```
>> bplot(sort.time, log10(sort.size))
```



This bplot shows us the relationship between input size and the resulting time. This graph shows us that there is a much larger range between times when the input size is large (4) than when it is smaller (2 and 3). Similar to the previous bplot, only 4 shows the actual box; 2 and 3 are skewed into a line and hard to read. We can, again, infer a relatively large standard deviation in 4 when compared to 2 and 3.

```
>> bplot(sort.time, sort.algorithm)
```



This bplot shows the relationship between the chosen algorithm and the resulting time. MergeSort and Quicksort had relatively tiny runtimes (always under 0.25 seconds), whereas BubbleSort and SelectionSort had very high (and widespread) runtimes (up to over 19 seconds). This is exactly what the graph shows. We can see that BubbleSort had a very wide range, from almost 0 seconds to almost 20 seconds. SelectionSort is in a similar situation, ranging from nearly 0 seconds to almost 10 seconds. This is shown by the whiskers in each of those graphs. For the remaining two algorithms, MergeSort and QuickSort, the box/whiskers is nearly invisible and the data is almost skewed into a straight line at the bottom. This shows that, relative to the other two algorithms, the change in MergeSort and QuickSort's runtimes as input size changes is almost negligible.



# ANOVA

```
>> lm(sort)
```

Variable Names in Input Dataset

```
size
time
language
algorithm
```

Enter class or model statement or type ex for examples

```
lm>> class size language algorithm
```

Enter model statement

```
lm>> model time
=language+algorithm+log10(size)+language*algorithm+language*log10(size)+algorithm*log10(size)+language*algorithm*log10(size)
```

Sequential Sums of Squares ANOVA Table

| Source                             | df | SS      | MS      | F      | P-val    |
|------------------------------------|----|---------|---------|--------|----------|
| Language                           | 2  | 48.1618 | 24.0809 | 3.7314 | 0.054937 |
| Algorithm                          | 3  | 31.3453 | 10.4484 | 1.6190 | 0.236910 |
| log <sub>10</sub> (size)           | 1  | 37.4917 | 37.4917 | 5.8094 | 0.032897 |
| language*algorithm                 | 6  | 59.0133 | 9.8355  | 1.5240 | 0.251360 |
| language*log <sub>10</sub> (size)  | 2  | 70.8134 | 35.4067 | 5.4863 | 0.020315 |
| algorithm*log <sub>10</sub> (size) | 3  | 46.1552 | 15.3851 | 2.3839 | 0.120350 |
| lang*alg*log <sub>10</sub> (size)  | 6  | 86.8783 | 14.4797 | 2.2437 | 0.110050 |

|       |    |         |        |
|-------|----|---------|--------|
| Error | 12 | 77.4435 | 6.4536 |
|-------|----|---------|--------|

R-square 0.83065

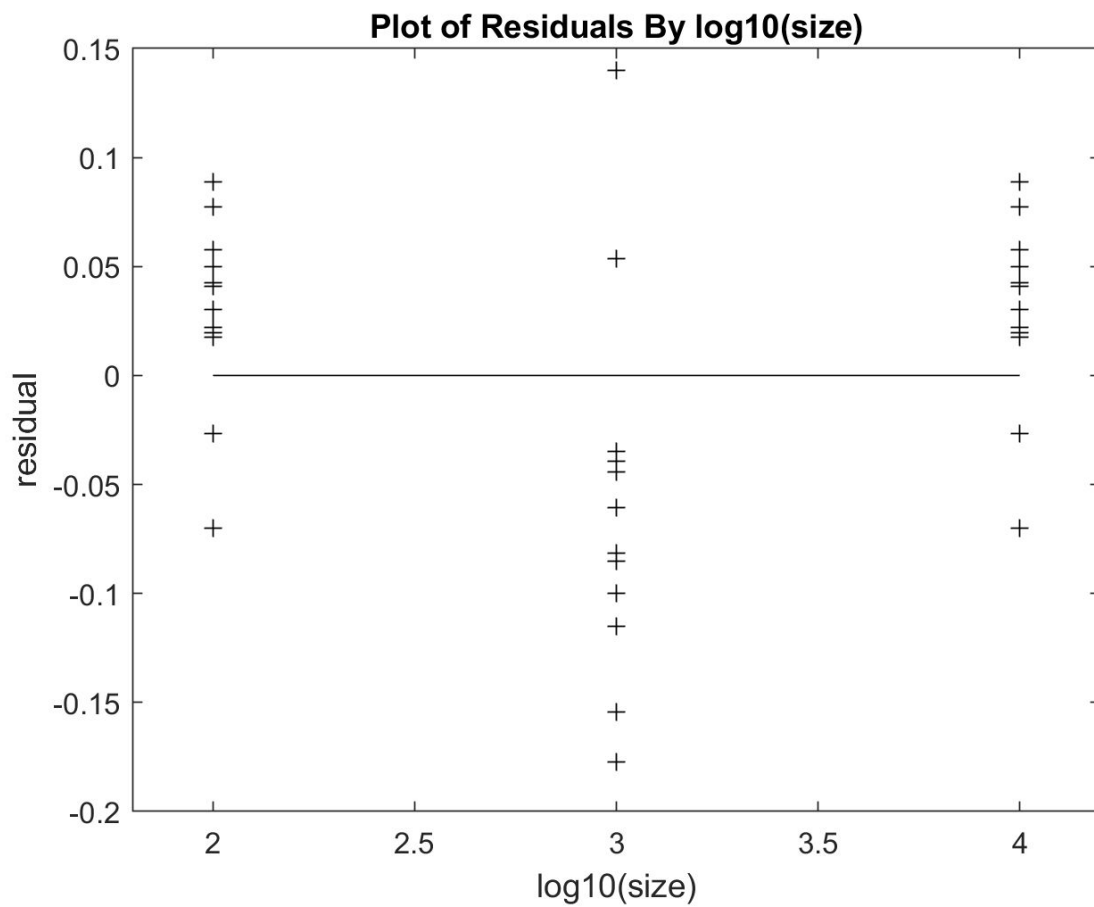
Standard Error 2.5404

```
lm>> model log(time)
=language+algorithm+log10(size)+language*algorithm+language*log10(size)+algorithm*log10(size)+language*algorithm*log10(size)
```

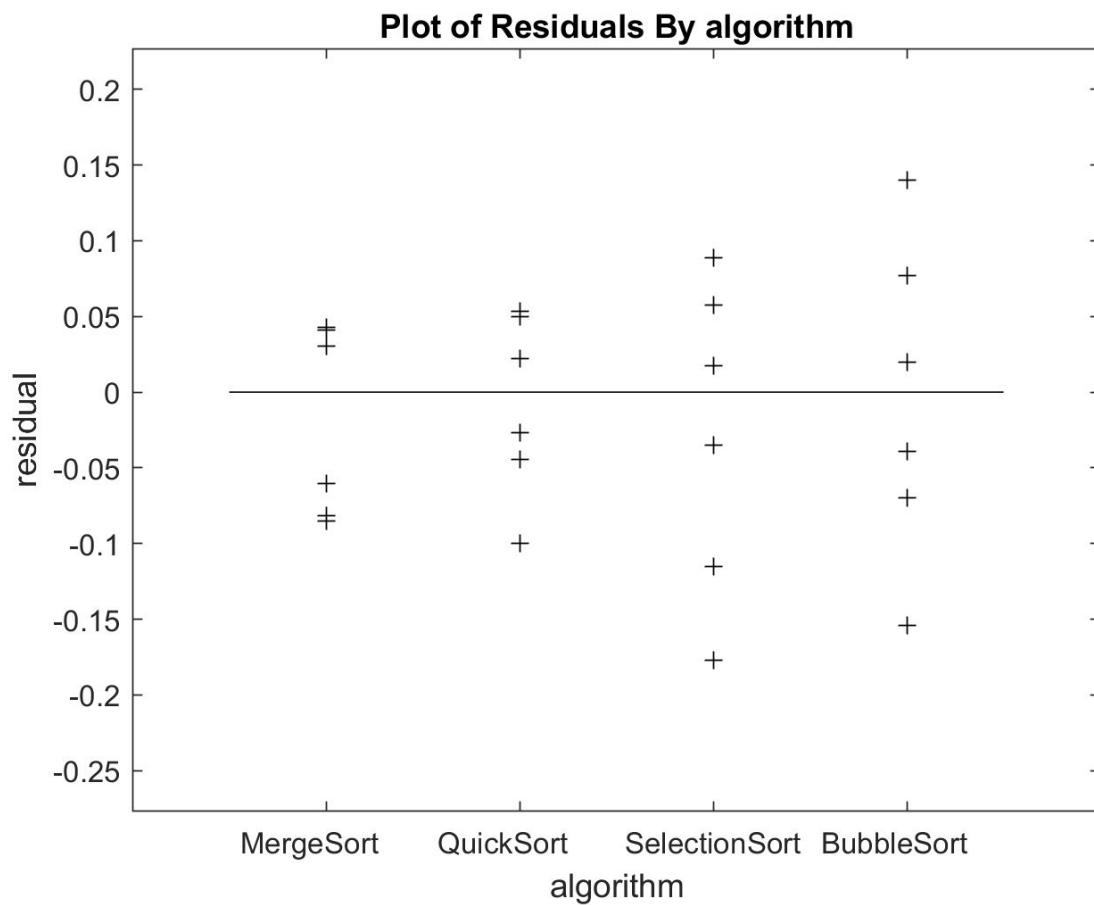
Sequential Sums of Squares ANOVA Table

| Source                              | df | SS        | MS         | F          | P-val      |
|-------------------------------------|----|-----------|------------|------------|------------|
| Language                            | 2  | 105.02210 | 52.511100  | 3410.7878  | 0.00000000 |
| Algorithm                           | 3  | 43.34800  | 14.449300  | 938.5374   | 1.7097e-14 |
| $\log_{10}(\text{size})$            | 1  | 197.71000 | 197.710000 | 12841.9947 | 0.00000000 |
| language*algorithm                  | 6  | 11.21130  | 1.868600   | 121.3693   | 4.9425e-10 |
| language* $\log_{10}(\text{size})$  | 2  | 2.01380   | 6.006900   | 390.1698   | 1.2068e-11 |
| algorithm* $\log_{10}(\text{size})$ | 3  | 32.68900  | 10.896300  | 707.7564   | 9.2149e-14 |
| lang*alg* $\log_{10}(\text{size})$  | 6  | 0.90689   | 0.151150   | 9.8177     | 0.00048115 |
| Error                               | 12 | 0.18475   | 0.015396   |            |            |
| R-square 0.99954                    |    |           |            |            |            |
| Standard Error 0.12408              |    |           |            |            |            |

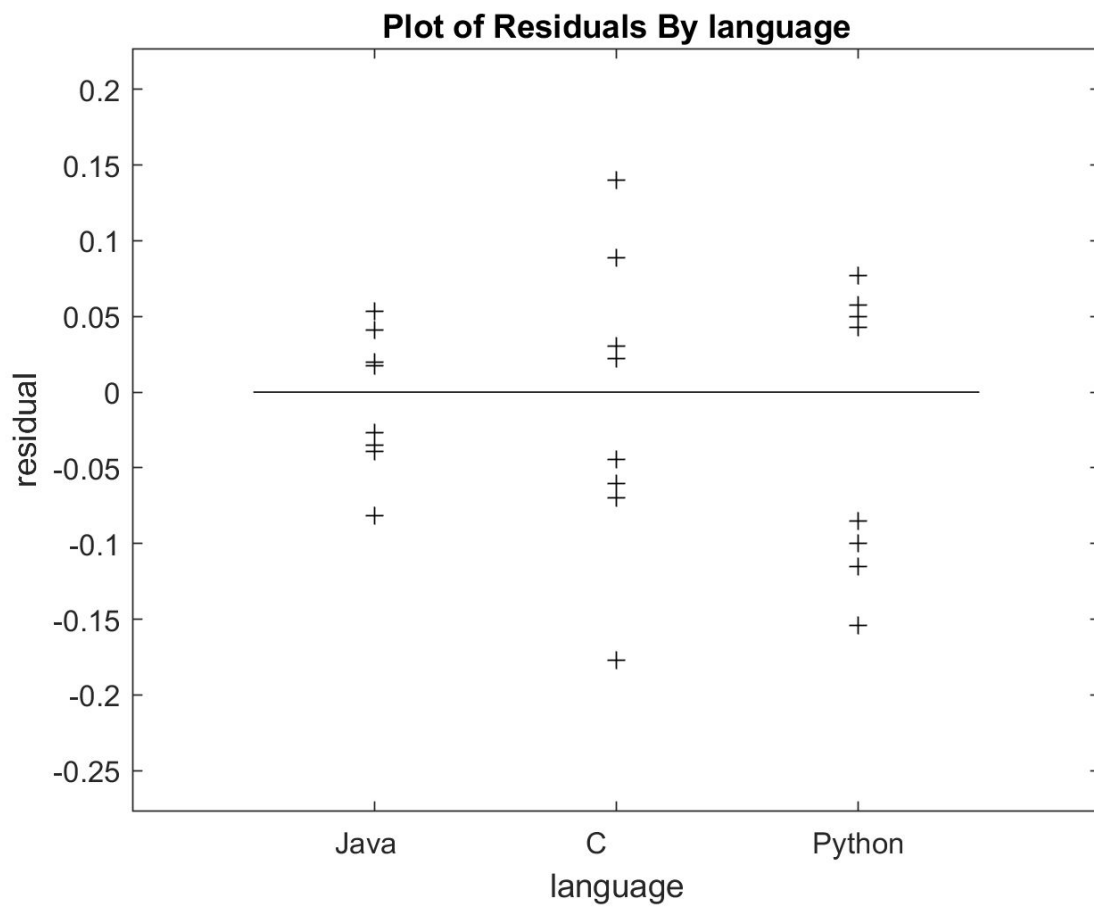
The ANOVA table above showed that  $\log_{10}(\text{size})$  and language\* $\log_{10}(\text{size})$  have statistical significance because they are the only two p-values which are less than 0.05. The R-square value is 0.83065; meaning that the graph is not as reliable as it could be in modeling the speed for sorting. The log model greatly increased the R-square value as well as decreasing the P-values so that all are below 0.05. The standard error was also decreased to .12408 from 2.5404. This indicates that the log model is much better at modeling the time it takes to sort data. Since the log model indicated a better fit, we can assume that the data shows a more exponential than a linear one.



The line of best fit from the lm model is compared to the actual data points to give us the residual plot. The plot may seem like it is not a very good fit because there are a lot of points below the line of best fit for size 3; however, the standard deviation is less than .2 throughout the plot. This tells us that the model produced by the lm function is a very good fit.



The residual plot for the algorithms is very similar to the residual plot of the size: it looks as if it indicates a bad fit with the model. The standard deviation on this plot does not exceed 0.2 as well, indicating a very good fit with the model that lm produced. Even in this residual plot, there is a much higher variation in bubblesort than there is in other sorting algorithms.



The residual plot for Language is similar throughout all the levels indicating that the log model fits the data very well. Even though the points look spread out, the labels on the axis show that the data is less than .2 standard deviations away from the model. The data is very close to the model that the lm function produced.

## Practical Implications

As mentioned before, the practical implications of this data are massive. Nearly every program, whether it seems like it at first or not, needs to sort a list at some point or another. It may be a list of users sorted alphabetically, a list of search results sorted by relevance, a list of locations sorted by distance... the list goes on and on. For major companies running software that is distributed globally, this knowledge is critical so that every user has a good experience, and there are minimal load times throughout.

I think a good take-home lesson from this study for the general population would be this: never use BubbleSort - (link [and if you still need more convincing, Obama would say the same thing](#))<sup>1</sup>. Beyond this, the lesson would be that regardless of the programming language you are using, if you have the time required to implement it, use QuickSort for your sorting needs. QuickSort proved to be the fastest across languages, and the sorting algorithm is almost always something you can choose -- frequently, in the middle of a project, the language is not. However, should language be up to you, and should sorting efficiency be a primary deciding factor, Java may be the way to go, as it was consistently faster across all tests.

---

<sup>1</sup> <https://www.youtube.com/watch?v=koMpGeZpu4Q>

## Further Questions

One of the primary questions that came to me during this study was how would these algorithms and languages compare on different machines? The actual times would likely be different, but would the ratios and models of the times be the same?

Another piece that isn't so much a question, but maybe another layer to the study, is how closely do these algorithms follow their proven, theoretical asymptotic runtimes? In practice, do they follow the expected trendline?

Lastly, this question could be asked - in terms of speed vs. input size, is this as efficient as it gets? Are there still more algorithms waiting to be developed that could bring runtime from  $O(n^2)$  and  $O(n \log n)$  to even  $O(\log n)$  or  $O(n)$ ?

If we were to do this project again, knowing what we know now, I think it would be interesting to go more in-depth on the algorithms themselves, and not use language as a factor. This would allow us to do more testing on more input sizes, and not worry about whether the language was influencing our results. We could also develop a model that would predict runtime given an input size, without worry of a language, if we calculated the average per-unit runtime in terms of  $n$  for sorting a list of size  $n$ .